

Decision Tree.

In this part, we use a classification algorithm to build a model from the historical data of prices. We then use the trained decision tree to predict the price of a future house sale.

First, we import the Following Libraries:

- **numpy** (as **np**)
- **pandas**
- **DecisionTreeClassifier** from **sklearn.tree**

```
In [2]: import numpy as np
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
```

```
In [3]: # Importing the dataset

df = pd.read_csv("nyc-rolling-sales.csv", delimiter=",")
df.head()
```

Out[3]:

	Unnamed: 0	BOROUGH	NEIGHBORHOOD	BUILDING CLASS CATEGORY	TAX CLASS AT PRESENT	BLOCK	LOT	EASE- MENT	BUI CL/ PRI
0	4	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2A	392	6		
1	5	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	399	26		
2	6	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2	399	39		
3	7	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	402	21		
4	8	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2A	404	55		

5 rows × 22 columns

Data processing & Exploration

```
In [4]: #Removing unnecessary columns
del df['EASE-MENT']
del df['Unnamed: 0']
del df['ADDRESS']
del df['APARTMENT NUMBER']
```

```
In [5]: #Checking for duplicates
sum(df.duplicated(df.columns))
```

Out[5]: 956

```
In [6]: #Removing duplicate records
df = df.drop_duplicates(df.columns, keep='last')
sum(df.duplicated(df.columns))
```

Out[6]: 0

```
In [7]: #Convert some of the columns to desired datatype
df['TAX CLASS AT TIME OF SALE'] = df['TAX CLASS AT TIME OF SALE'].astype(
    'category')
df['TAX CLASS AT PRESENT'] = df['TAX CLASS AT PRESENT'].astype('category')
df['LAND SQUARE FEET'] = pd.to_numeric(df['LAND SQUARE FEET'], errors='coerce')
df['GROSS SQUARE FEET'] = pd.to_numeric(df['GROSS SQUARE FEET'], errors='coerce')
df['SALE PRICE'] = pd.to_numeric(df['SALE PRICE'], errors='coerce')
df['BOROUGH'] = df['BOROUGH'].astype('category')
```

```
In [8]: # Convert Sale Date to Year
from datetime import datetime

for i in range(len(df)):
    if True:
        the_date = datetime.strptime(str(df['SALE DATE'][i]), '%Y-%m-%d
%H:%M:%S')
        df.at[i, 'SALE DATE'] = the_date.year
    else:
        df.at[i, 'SALE DATE'] = int(df.at[i, 'SALE DATE'])

# convert to integer
df['SALE DATE'] = df['SALE DATE'].astype(int)

df['SALE DATE'].head()
```

```

-----
-----
KeyError                                Traceback (most recent call last)
<ipython-input-8-c8db02fad33a> in <module>
      4 for i in range(len(df)):
      5     if True:
----> 6         the_date = datetime.strptime(str(df['SALE DATE'][i]),
      7         '%Y-%m-%d %H:%M:%S')
      8         df.at[i, 'SALE DATE'] = the_date.year
      9     else:

//anaconda3/lib/python3.7/site-packages/pandas/core/series.py in __getitem__(self, key)
    866         key = com.apply_if_callable(key, self)
    867         try:
--> 868             result = self.index.get_value(self, key)
    869
    870             if not is_scalar(result):

//anaconda3/lib/python3.7/site-packages/pandas/core/indexes/base.py in get_value(self, series, key)
    4373         try:
    4374             return self._engine.get_value(s, k,
-> 4375                                         tz=getattr(series.dtype,
pe, 'tz', None))
    4376         except KeyError as e1:
    4377             if len(self) > 0 and (self.holds_integer() or self.is_boolean()):

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_value()

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_value()

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.Int64HashTable.get_item()

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.Int64HashTable.get_item()

KeyError: 48

```

```

In [10]: #checking missing values
df.columns[df.isnull().any()]

```

```

Out[10]: Index(['LAND SQUARE FEET', 'GROSS SQUARE FEET', 'SALE PRICE'], dtype='object')

```

```
In [11]: miss=df.isnull().sum()/len(df)
miss=miss[miss>0]
miss.sort_values(inplace=True)
miss
```

```
Out[11]: SALE PRICE          0.168365
LAND SQUARE FEET        0.310484
GROSS SQUARE FEET       0.326371
dtype: float64
```

```
In [12]: #Convert series to column DataFrame
miss=miss.to_frame()
#Set Column Name
miss.columns=['count']
#Set Index Name
miss.index.names=['Name']
#Create Column from Index
miss['Name']=miss.index
miss
```

```
Out[12]:
```

	count	Name
SALE PRICE	0.168365	SALE PRICE
LAND SQUARE FEET	0.310484	LAND SQUARE FEET
GROSS SQUARE FEET	0.326371	GROSS SQUARE FEET

```
In [13]: #Plot the missing values
import seaborn as sns
import matplotlib.pyplot as plt

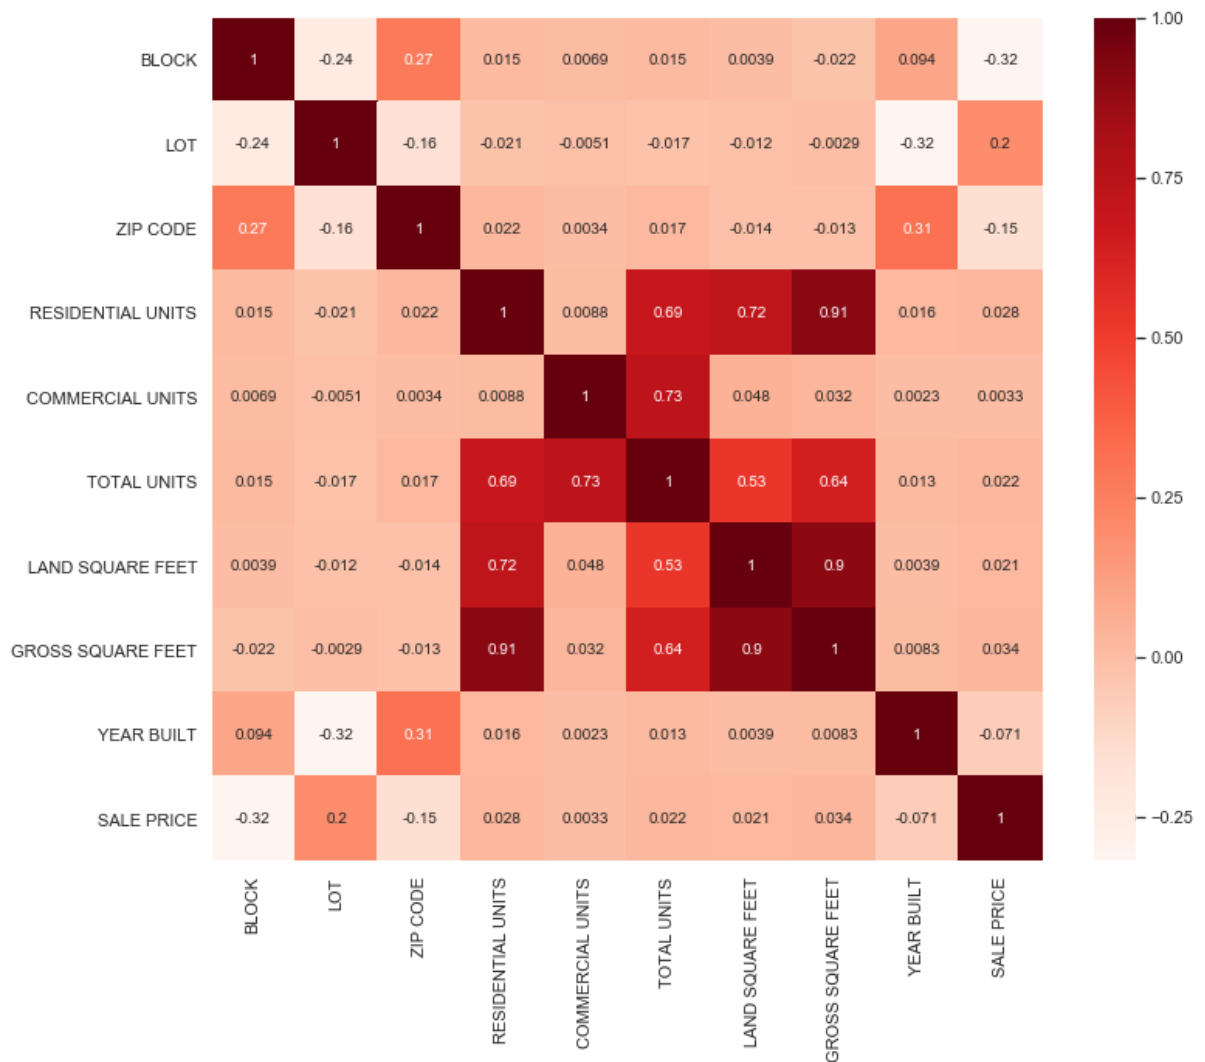
sns.set(style='whitegrid',color_codes=True)
sns.barplot(x='Name', y='count',data=miss)
plt.xticks(rotation=90)
sns
```

```
Out[13]: <module 'seaborn' from '//anaconda3/lib/python3.7/site-packages/seabor
n/__init__.py'>
```

```
In [14]: #Populating mean values for missing data
df['LAND SQUARE FEET']=df['LAND SQUARE FEET'].fillna(df['LAND SQUARE FEE
T'].mean())
df['GROSS SQUARE FEET']=df['GROSS SQUARE FEET'].fillna(df['GROSS SQUARE
FEET'].mean())
```

```
In [16]: # Removing null observations
df = df[(df['SALE PRICE'] > 100000) & (df['SALE PRICE'] < 5000000)]
```

```
In [19]: #Using Pearson Correlation
plt.figure(figsize=(12,10))
cor = df.corr()
sns.heatmap(cor, annot=True, cmap=plt.cm.Red)
plt.show()
```



```
In [20]: # Most Correlations
cor_target = abs(cor["SALE PRICE"])
#Selecting highly correlated features
relevant_features = cor_target[cor_target>0.1]
relevant_features
```

```
Out[20]: BLOCK      0.319305
LOT      0.195200
ZIP CODE  0.151620
SALE PRICE 1.000000
Name: SALE PRICE, dtype: float64
```

```
In [22]: del df['SALE DATE']
```

Encoding

```
In [24]: X = df[['BOROUGH', 'NEIGHBORHOOD', 'BUILDING CLASS CATEGORY', 'TAX CLASS AT  
PRESENT', 'BLOCK', 'LOT', 'BUILDING CLASS AT PRESENT', 'ZIP CODE', 'RESIDENTI  
AL UNITS', 'COMMERCIAL UNITS', 'TOTAL UNITS', 'LAND SQUARE FEET', 'GROSS SQU  
ARE FEET', 'YEAR BUILT', 'TAX CLASS AT TIME OF SALE', 'BUILDING CLASS AT TI  
ME OF SALE']].values  
X[:, 14]
```

```
Out[24]: array([2, 2, 2, ..., 1, 1, 1], dtype=object)
```

```
In [25]: # Getting the dependent variables and independent variables  
X = df[['BOROUGH', 'NEIGHBORHOOD', 'BUILDING CLASS CATEGORY', 'TAX CLASS AT  
PRESENT', 'BLOCK', 'LOT', 'BUILDING CLASS AT PRESENT', 'ZIP CODE', 'RESIDENTI  
AL UNITS', 'COMMERCIAL UNITS', 'TOTAL UNITS', 'LAND SQUARE FEET', 'GROSS SQU  
ARE FEET', 'YEAR BUILT', 'TAX CLASS AT TIME OF SALE', 'BUILDING CLASS AT TI  
ME OF SALE']].values  
  
# Encoding categorical data  
from sklearn.preprocessing import LabelEncoder, OneHotEncoder  
labelencoder_X_1 = LabelEncoder()  
X[:, 1] = labelencoder_X_1.fit_transform(X[:, 1])  
  
labelencoder_X_2 = LabelEncoder()  
X[:, 2] = labelencoder_X_2.fit_transform(X[:, 2])  
  
labelencoder_X_3 = LabelEncoder()  
X[:, 3] = labelencoder_X_3.fit_transform(X[:, 3])  
  
labelencoder_X_6 = LabelEncoder()  
X[:, 6] = labelencoder_X_6.fit_transform(X[:, 6])  
  
labelencoder_X_16 = LabelEncoder()  
X[:, 15] = labelencoder_X_16.fit_transform(X[:, 15])
```

```
In [26]: X[0:5]
```

```
Out[26]: array([[1, 1, 6, 7, 402, 21, 18, 10009, 10, 0, 10, 2272.0, 6794.0, 191  
3,  
2, 17],  
[1, 1, 6, 7, 406, 32, 18, 10009, 8, 0, 8, 1750.0, 4226.0, 1920,  
2,  
17],  
[1, 1, 8, 5, 373, 40, 20, 10009, 0, 0, 0, 3846.981435858288,  
3874.3228378618364, 1920, 2, 19],  
[1, 1, 8, 5, 373, 40, 20, 10009, 0, 0, 0, 3846.981435858288,  
3874.3228378618364, 1920, 2, 19],  
[1, 1, 8, 5, 373, 40, 20, 10009, 0, 0, 0, 3846.981435858288,  
3874.3228378618364, 1920, 2, 19]], dtype=object)
```

We then set the target variable, Sale Price

```
In [28]: # Target variable, Sale Price

y = df['SALE PRICE'].values
y[0:5]
```

```
Out[28]: array([3936272., 3192840., 499000., 529500., 423000.])
```

Setting up the decision tree

```
In [29]: import numpy as np
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
```

```
In [30]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

Now **train_test_split** will return 4 different parameters. We will name them:

X_train, X_test, y_train, y_test

The **train_test_split** will need the parameters:

X, y, test_size=0.3, and random_state=34.

We will also try random state = 0

The **X** and **y** are the arrays required before the split, the **test_size** represents the ratio of the testing dataset, and the **random_state** ensures that we obtain the same splits.

```
In [31]: # Splitting the training set and test set
X_train ,X_test, y_train , y_test = train_test_split(X , y, test_size =
0.3 , random_state =34)
```

```
In [32]: # Training set
X_train.shape , y_train.shape
```

```
Out[32]: ((38173, 16), (38173,))
```

```
In [33]: # Feature Scaling
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```



```
In [34]: # Checking shape of X_train and y_train. Thier shapes should match
print(X_train.shape)
print(y_train.shape)

(38173, 16)
(38173,)
```

```
In [35]: # Also checking the shape of X_test and y_test. Thier shapes should also
match
print(X_test.shape)
print(y_test.shape)

(16361, 16)
(16361,)
```

Modeling

We will first create an instance of the **DecisionTreeClassifier** called **priceTree**.

Inside of the classifier, specify *criterion="entropy"* so we can see the information gain of each node.

```
In [36]: priceTree = DecisionTreeClassifier(criterion="entropy", max_depth = 9)
priceTree # it shows the default parameters
```

```
Out[36]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_dept
h=9,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=No
ne,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False,
                                random_state=None, splitter='best')
```

Next, we will fit the data with the training feature matrix **X_trainset** and training response vector **y_trainset**

```
In [37]: priceTree.fit(X_train, y_train)
```

```
Out[37]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_dept
h=9,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=No
ne,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False,
                                random_state=None, splitter='best')
```

Prediction

Let's make some **predictions** on the testing dataset and store it into a variable called **predTree**.

```
In [38]: predTree = priceTree.predict(X_test)
```

We print out **predTree** and **y_test** to visually compare the prediction to the actual values.

```
In [39]: print (predTree [0:5])
print (y_test [0:5])

[420000. 400000. 763687. 470000. 150000.]
[1270000. 2500000. 491790. 332000. 140000.]
```

Evaluation

Next, let's import **metrics** from sklearn and check the accuracy of our model.

```
In [40]: from sklearn import metrics
import matplotlib.pyplot as plt
print("DecisionTrees's Accuracy: ", metrics.accuracy_score(y_test, predT
ree))
```

```
DecisionTrees's Accuracy: 0.015035755760650327
```

```
In [73]: # RMSE
from sklearn.metrics import mean_squared_error
def rmse(y_test,y_pred):
    return np.sqrt(mean_squared_error(y_test,y_pred))
```

```
In [75]: rmse(y_test,predTree)
```

```
Out[75]: 692152.7732118465
```

These are very poor scores. We will do some visulization, check feature importance and re-build the model

Feature Importance

With the very low accuracy, we would like to review the data and remove any noise that may be affecting performance

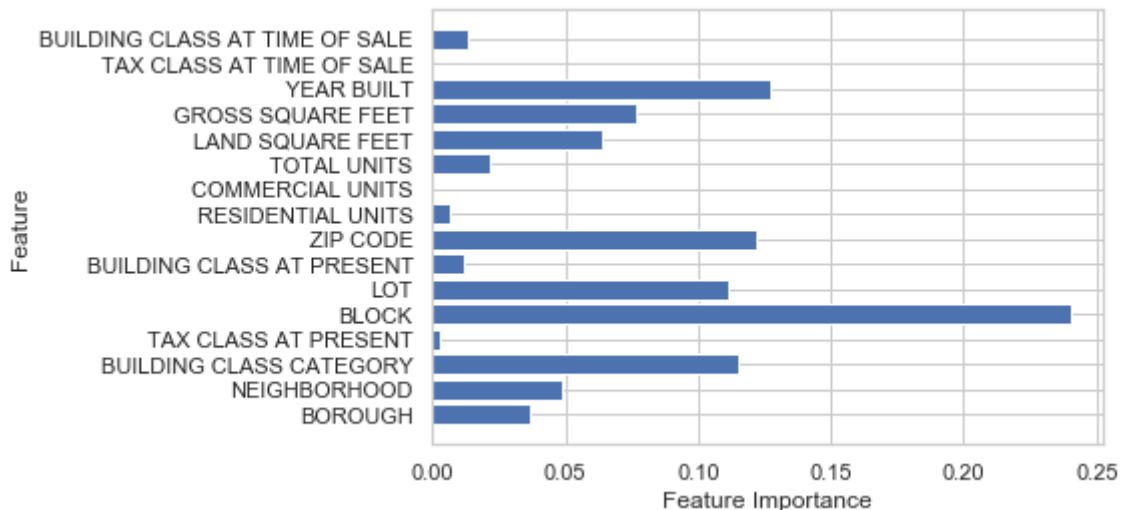
```
In [41]: from sklearn.ensemble import ExtraTreesClassifier
```

```
In [42]: print('Feature importances: {}'.format(priceTree.feature_importances_))
         type(priceTree.feature_importances_)
```

```
Feature importances: [0.03679404 0.04857123 0.11549536 0.00308566 0.240
43385 0.11155583
0.01218395 0.12234182 0.00657886 0.          0.02139389 0.06378611
0.07683286 0.12721396 0.          0.01373256]
```

Out[42]: numpy.ndarray

```
In [43]: #Feature Importance
n_feature = X_train.shape[1]
featureNames = df.drop(columns='SALE PRICE')
featureNames = featureNames.columns.values
plt.barh(range(n_feature), priceTree.feature_importances_, align='center')
plt.yticks(np.arange(n_feature), featureNames)
plt.xlabel('Feature Importance')
plt.ylabel('Feature')
plt.show()
```



Model Re-building

We will now re-train the data using selected features

- Features with importance of less than .05 will be dropped
- We can also see that the dropped features do not correlate well with SALE PRICE
- Finally, we want to check that the top important features do not correlate much with each other

```
In [44]: # Most Correlations
cor_target2 = abs(cor["SALE PRICE"])
#Selecting highly correlated features
relevant_features2 = cor_target2[cor_target2>0.1]
relevant_features2
```

```
Out[44]: BLOCK          0.319305
LOT          0.195200
ZIP CODE     0.151620
SALE PRICE   1.000000
Name: SALE PRICE, dtype: float64
```

```
In [45]: df.head()
```

```
Out[45]:
```

	BOROUGH	NEIGHBORHOOD	BUILDING CLASS CATEGORY	TAX CLASS AT PRESENT	BLOCK	LOT	BUILDING CLASS AT PRESENT	ZIP CODE	RE
3	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	402	21	C4	10009	
6	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	406	32	C4	10009	
13	1	ALPHABET CITY	09 COOPS - WALKUP APARTMENTS	2	373	40	C6	10009	
15	1	ALPHABET CITY	09 COOPS - WALKUP APARTMENTS	2	373	40	C6	10009	
16	1	ALPHABET CITY	09 COOPS - WALKUP APARTMENTS	2	373	40	C6	10009	

```
In [46]: df_2 = []
df_2 = df
df_2.head()
```

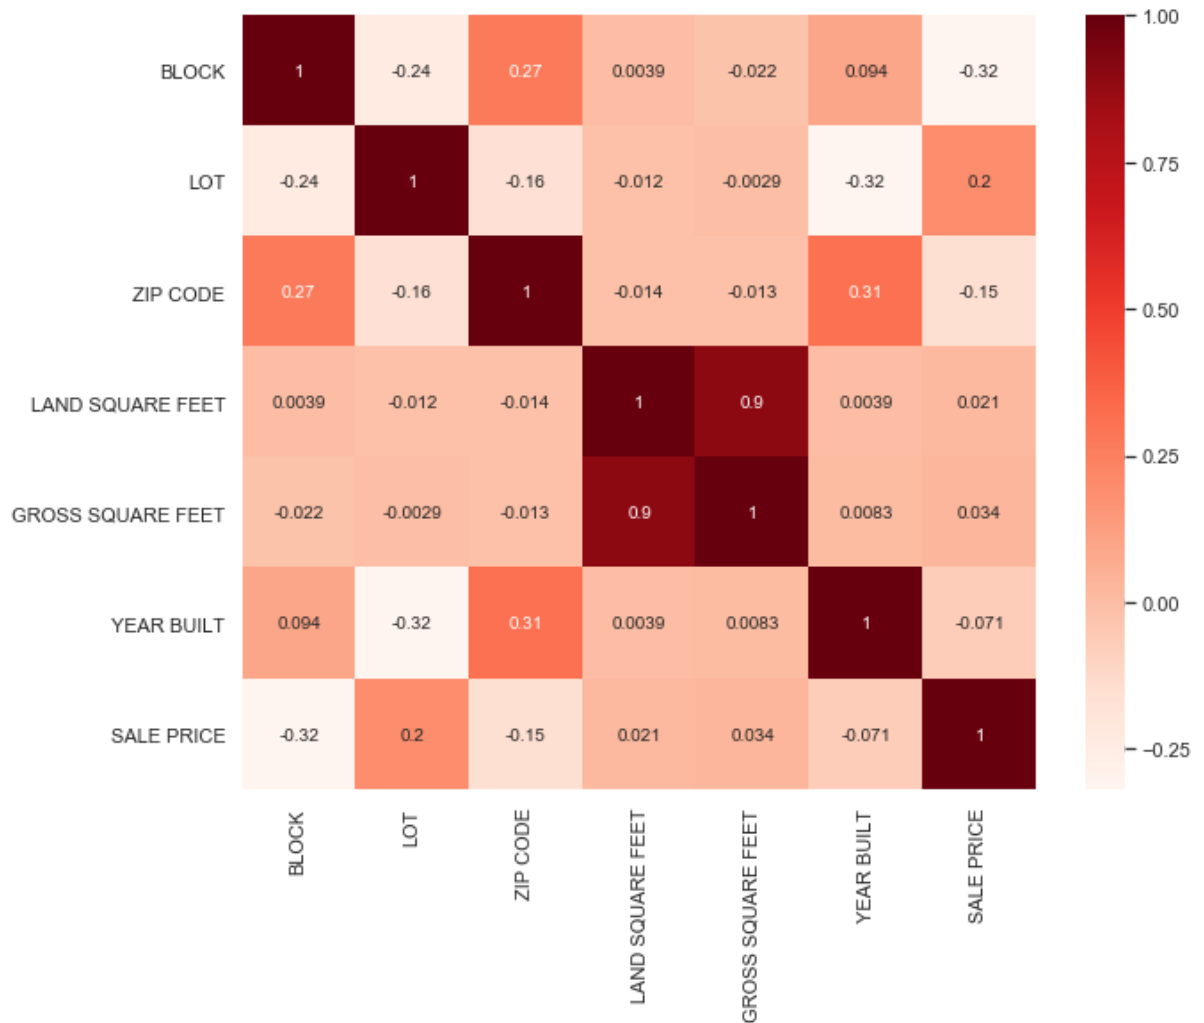
Out[46]:

	BOROUGH	NEIGHBORHOOD	BUILDING CLASS CATEGORY	TAX CLASS AT PRESENT	BLOCK	LOT	BUILDING CLASS AT PRESENT	ZIP CODE	RE
3	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	402	21	C4	10009	
6	1	ALPHABET CITY	07 RENTALS - WALKUP APARTMENTS	2B	406	32	C4	10009	
13	1	ALPHABET CITY	09 COOPS - WALKUP APARTMENTS	2	373	40	C6	10009	
15	1	ALPHABET CITY	09 COOPS - WALKUP APARTMENTS	2	373	40	C6	10009	
16	1	ALPHABET CITY	09 COOPS - WALKUP APARTMENTS	2	373	40	C6	10009	

```
In [47]: # remove the less important features

del df_2['BUILDING CLASS AT TIME OF SALE']
del df_2['TAX CLASS AT TIME OF SALE']
del df_2['TOTAL UNITS']
del df_2['COMMERCIAL UNITS']
del df_2['RESIDENTIAL UNITS']
del df_2['BUILDING CLASS AT PRESENT']
del df_2['TAX CLASS AT PRESENT']
del df_2['NEIGHBORHOOD']
del df_2['BOROUGH']
# del df_2['SALE PRICE'] # do not delete sale data yet so we can run co
rrelation
```

```
In [49]: #Using Pearson Correlation to check feature correlations
plt.figure(figsize=(10,8))
cor2 = df_2.corr()
sns.heatmap(cor2, annot=True, cmap=plt.cm.Reds)
plt.show()
```



```
In [50]: # GROSS SQUARE FEET and LAND SQUARE FEET are highly correlated with each other,
# We drop LAND SQUARE FEET because it is slightly less correlated with SALE PRICE at .021

del df_2['LAND SQUARE FEET']
```

We now begin rebuilding the model

```
In [51]: df_2.head()
```

Out[51]:

	BUILDING CLASS CATEGORY	BLOCK	LOT	ZIP CODE	GROSS SQUARE FEET	YEAR BUILT	SALE PRICE
3	07 RENTALS - WALKUP APARTMENTS	402	21	10009	6794.000000	1913	3936272.0
6	07 RENTALS - WALKUP APARTMENTS	406	32	10009	4226.000000	1920	3192840.0
13	09 COOPS - WALKUP APARTMENTS	373	40	10009	3874.322838	1920	499000.0
15	09 COOPS - WALKUP APARTMENTS	373	40	10009	3874.322838	1920	529500.0
16	09 COOPS - WALKUP APARTMENTS	373	40	10009	3874.322838	1920	423000.0

```
In [52]: df_2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 54534 entries, 3 to 84545
Data columns (total 7 columns):
BUILDING CLASS CATEGORY    54534 non-null object
BLOCK                      54534 non-null int64
LOT                        54534 non-null int64
ZIP CODE                   54534 non-null int64
GROSS SQUARE FEET         54534 non-null float64
YEAR BUILT                 54534 non-null int64
SALE PRICE                 54534 non-null float64
dtypes: float64(2), int64(4), object(1)
memory usage: 3.3+ MB
```

```
In [53]: # Get the variables
A = df_2[['BUILDING CLASS CATEGORY', 'BLOCK', 'LOT', 'ZIP CODE', 'GROSS SQUARE FEET', 'YEAR BUILT']].values
b = df_2['SALE PRICE'].values

# Encoding categorical data
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
labelencoder_A_1 = LabelEncoder()
A[:, 0] = labelencoder_A_1.fit_transform(A[:, 0])

# Target variable, Sale Price

A[0:5]
b[0:5]
```

Out[53]: array([3936272., 3192840., 499000., 529500., 423000.])

```
In [54]: # Split training set and test set
A_train, A_test, b_train, b_test = train_test_split(A, b, test_size = 0.3, random_state = 34)
```

```
In [55]: # Training set
A_train.shape , b_train.shape
```

```
Out[55]: ((38173, 6), (38173,))
```

```
In [56]: # Feature Scaling
sc = StandardScaler()
A_train = sc.fit_transform(A_train)
A_test = sc.transform(A_test)
```

```
In [57]: print(A_train.shape)
print(b_train.shape)

print(A_test.shape)
print(b_test.shape)
```

```
(38173, 6)
(38173,)
(16361, 6)
(16361,)
```

```
In [58]: # 1 specify second model
priceTree_2 = DecisionTreeClassifier(criterion="entropy", max_depth = 9)
priceTree_2 # it shows the default parameters
```

```
Out[58]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_dept
h=9,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=No
ne,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False,
                                random_state=None, splitter='best')
```

```
In [59]: # 2 fit the data
priceTree_2.fit(A_train, b_train)
```

```
Out[59]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_dept
h=9,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=No
ne,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False,
                                random_state=None, splitter='best')
```

```
In [60]: # 3 Prediction
predTree_2 = priceTree_2.predict(A_test)
```



```
In [61]: # 4 Print prediction
print (predTree_2 [0:5])
print (b_test [0:5])

[1185000.  400000. 1050000.  470000.  250000.]
[1270000. 2500000.  491790.  332000.  140000.]
```

```
In [62]: # 5 Accuracy of predTree_2
print("DecisionTrees's Accuracy: ", metrics.accuracy_score(b_test, predTree_2))
```

DecisionTrees's Accuracy: 0.0149746348022737

After feature selction, the Accuracy is relatively the same!

We will build one more model using a higher max_depth

```
In [63]: # Model using max_depth = 20

priceTree_3 = DecisionTreeClassifier(criterion="entropy", max_depth = 20
)
priceTree_3.fit(A_train, b_train)
predTree_3 = priceTree_3.predict(A_test)

print("DecisionTrees's Accuracy: ", metrics.accuracy_score(b_test, predTree_3))
```

DecisionTrees's Accuracy: 0.01772507792922193

F1 score, Jaccard Similarity, and RMSE

```
In [64]: from sklearn.metrics import f1_score
f1_score(b_test, predTree_3, average='weighted')

//anaconda3/lib/python3.7/site-packages/sklearn/metrics/classification.py:1437: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no predicted samples.
'precision', 'predicted', average, warn_for)
//anaconda3/lib/python3.7/site-packages/sklearn/metrics/classification.py:1439: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no true samples.
'recall', 'true', average, warn_for)
```

Out[64]: 0.016880589784208252

```
In [68]: from sklearn.metrics import jaccard_similarity_score
jaccard_similarity_score(b_test, predTree_3)
```

```
//anaconda3/lib/python3.7/site-packages/sklearn/metrics/classification.py:635: DeprecationWarning: jaccard_similarity_score has been deprecated and replaced with jaccard_score. It will be removed in version 0.23. This implementation has surprising behavior for binary and multiclass classification tasks.
'and multiclass classification tasks.', DeprecationWarning)
```

```
Out[68]: 0.01772507792922193
```

```
In [76]: rmse(b_test, predTree_3)
```

```
Out[76]: 620436.6733025813
```

Conclusion

With a higher max_depth, and with feature selection, the Accuracy is still very low. It appears that Decision Trees is not a great option for our problem

The Accuracy may be improved by grouping the price ranges into bins as this is what Decision trees would be great for